

# P1 Presentation

By: Alan Liang





# Overview of Process

1. Insert filtered data in table
2. Combine tables with filtered data (aggregate data if need to)
3. Do necessary computations in the table with filtered data
4. Combine results (columns) into one single table
5. Display the results with limit or where



# 1. Which English Wikipedia article got the most traffic on October 20th?

## DATA USED

pageviews-20201020-000000.gz -

pageviews-20201020-230000.gz

## HIVE QUERIES USED TO ANSWER THE QUESTION

First we set up our database and create a table that takes in pageview data from Oct containing only domain codes 'en' and 'en.m'.

1) CREATE DATABASE PAGEVIEW\_DB;

2) USE PAGEVIEW\_DB;

3) CREATE TABLE OCT\_PAGEVIEWS (DOMAIN\_CODE STRING, PAGE\_TITLE STRING, NUM\_OF\_REQUESTS INT, RESPONSE\_BYTES INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';

4) CREATE TABLE EN\_PAGEVIEWS (PAGE\_TITLE STRING, NUM\_OF\_REQUESTS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';

(cont in next page...)

# 1. Which English Wikipedia article got the most traffic on October 20th?

## HIVE QUERIES USED TO ANSWER THE QUESTION (cont.)

5) LOAD DATA LOCAL INPATH '/home/aliang30/pageviews-20201020/\*' INTO TABLE OCT\_PAGEVIEWS;

6) INSERT INTO EN\_PAGEVIEWS (SELECT PAGE\_TITLE, NUM\_OF\_REQUESTS FROM OCT\_PAGEVIEWS WHERE DOMAIN\_CODE='en');

7) INSERT INTO EN\_PAGEVIEWS (SELECT PAGE\_TITLE, NUM\_OF\_REQUESTS FROM OCT\_PAGEVIEWS WHERE DOMAIN\_CODE='en.m');


(Now that we have our table with only English wiki articles, we want to aggregate the num\_of\_requests with similar page\_titles.

And then we display our results

```
8) SELECT PAGE_TITLE,  
SUM (NUM_OF_REQUESTS) AS num_of_page_views FROM EN_PAGEVIEWS  
GROUP BY PAGE_TITLE  
ORDER BY SUM(NUM_OF_REQUESTS) DESC  
LIMIT 10;
```

Total MapReduce CPU Time Spent: 10 minutes 6 seconds  
OK

page_title	num_of_page_views
Main_Page	5961008
Special:Search	1476831
-	544714
Jeffrey_Toobin	321459
C._Rajagopalachari	210558
The_Haunting_of_Bly_Manor	185139
Robert_Redford	178779
Jeff_Bridges	159163
Bible	151484
Chicago_Seven	149966



## 2. What English Wikipedia article has the largest fraction of its readers follow an internal link to another wikipedia article?

### DATA USED

We will be using September's data

pageviews-20200901-000000.gz - pageviews-20200930-230000.gz

clickstream-enwiki-2020-09.tsv

### COMPUTATION USED

Percentage of readers follow an internal link =  $(\text{Total internal\_link clicks in article}) * 100 / \text{Total pageviews}$



## 2. What English Wikipedia article has the largest fraction of its readers follow an internal link to another wikipedia article?

### HIVE QUERIES USED TO ANSWER THE QUESTION

First we set up our database and create a table that takes in clickstream data from Sept containing only internal link type.

1) CREATE DATABASE CLICKSTREAM\_DB;

2) USE CLICKSTREAM\_DB;

3) CREATE TABLE CS\_TABLE (PREV STRING, CURR STRING, TYPE STRING, NUM\_OF\_CLICK\_THROUGHS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';

4) CREATE TABLE INTERNAL\_DATA (PREV STRING, NUM\_OF\_CLICK\_THROUGHS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';

5) LOAD DATA LOCAL INPATH '/home/aliang30/clickstream-enwiki-2020-09.tsv' INTO TABLE CS\_TABLE;

6) INSERT INTO INTERNAL\_DATA (SELECT PREV, NUM\_OF\_CLICK\_THROUGHS FROM CS\_TABLE WHERE TYPE='link');

Now that we have the data that contains only internal links, we must aggregate the num of internal link clicks with similar page\_titles and insert new data into a new table

7) CREATE TABLE UPDATED\_INTERNAL\_DATA (PREV STRING, NUM\_OF\_CLICK\_THROUGHS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';

8) INSERT INTO UPDATED\_INTERNAL\_DATA (SELECT PREV, SUM ( NUM\_OF\_CLICK\_THROUGHS) AS NUM\_OF\_INTERNAL FROM EN\_PAGEVIEWS GROUP BY PAGE\_TITLE);



## 2. What English Wikipedia article has the largest fraction of its readers follow an internal link to another wikipedia article?

### HIVE QUERIES USED TO ANSWER THE QUESTION (cont))

Now we need to create a new table that takes in pageview data from Sept containing only domain codes 'en' and 'en.m'.

```
7) CREATE TABLE SEPT_PAGEVIEWS (DOMAIN_CODE STRING, PAGE_TITLE STRING, NUM_OF_REQUESTS INT, RESPONSE_BYTES INT)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';
```

```
8) CREATE TABLE EN_PAGEVIEWS (PAGE_TITLE STRING, NUM_OF_REQUESTS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY '
';
```

```
9) LOAD DATA LOCAL INPATH '/home/aliang30/pageviews-20200920/*' INTO TABLE SEPT_PAGEVIEWS;
```

```
10) INSERT INTO EN_PAGEVIEWS (SELECT PAGE_TITLE, NUM_OF_REQUESTS FROM SEPT_PAGEVIEWS WHERE DOMAIN_CODE='en');
```

```
11) INSERT INTO EN_PAGEVIEWS (SELECT PAGE_TITLE, NUM_OF_REQUESTS FROM SEPT_PAGEVIEWS WHERE DOMAIN_CODE='en.m');
```

We aggregate the num\_of\_requests with similar page\_titles and insert new data into a separate table

```
12) CREATE TABLE UPDATED_EN_PAGEVIEWS (PAGE_TITLE STRING, NUM_OF_REQUESTS INT) ROW FORMAT DELIMITED FIELDS
TERMINATED BY ',';
```

```
13) INSERT INTO UPDATED_EN_PAGEVIEWS (SELECT PAGE_TITLE, SUM (NUM_OF_REQUESTS) AS NUM_OF_PAGE_VIEWS FROM
EN_PAGEVIEWS GROUP BY PAGE_TITLE);
```



## 2. What English Wikipedia article has the largest fraction of its readers follow an internal link to another wikipedia article?

### HIVE QUERIES USED TO ANSWER THE QUESTION (cont)

Now that we have our tables that contain 1) aggregated num\_of\_requests and 2) aggregated num\_of\_click\_throughs, we will use an inner join to return the rows that have matching page\_title.

14) CREATE TABLE Q2\_RESULT AS

```
SELECT UPDATED_EN_PAGEVIEWS.PAGE_TITLE, UPDATED_EN_PAGEVIEWS.NUM_OF_PAGE_VIEWS ,  
UPDATED_INTERNAL_DATA.NUM_OF_INTERNAL
```

```
FROM UPDATED_EN_PAGEVIEWS
```

```
INNER JOIN UPDATED_INTERNAL_DATA. ON UPDATED_EN_PAGEVIEWS.PAGE_TITLE = UPDATED_INTERNAL_DATA.PREV;
```

Now we want to do our computation:  $(\text{Total internal\_link clicks in article Link total} * 100 / \text{Total pageview})$  and store that in the PERCENTAGE\_OF\_READERS\_FOLLOW\_A\_LINK column.

15) SELECT PAGE\_TITLE, NUM\_OF\_PAGE\_VIEWS, NUM\_OF\_INTERNAL,


```
CAST(NUM_OF_INTERNAL * 100.0 / NUM_OF_PAGE_VIEWS AS DECIMAL) AS PERCENTAGE_OF_READERS_FOLLOW_A_LINK
```

```
FROM Q2_RESULT
```

```
ORDER BY PERCENTAGE_OF_READERS_FOLLOW_A_LINK DESC
```

```
LIMIT 10;
```





### 3. What series of Wikipedia articles, starting with Hotel California, keeps the largest fraction of its readers clicking on internal links?

#### HIVE QUERIES USED TO ANSWER THE QUESTION

Similar to question 2 but we will have to gather all of the articles "Hotel\_California" refers to and see which one has the highest number of internal click\_throughs

```
1) CREATE TABLE INTERNAL_DATA2 (PREV STRING, CURR STRING, NUM_OF_CLICK_THROUGHS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

We'll use the CS\_TABLE from previous question and gather all of the articles starting with "Hotel\_California" as well as type "link" and insert them in UPDATED\_INTERNAL\_DATA2.

```
2) INSERT INTO INTERNAL_DATA2
```

```
SELECT PREV, CURR, NUM_OF_CLICK_THROUGHS FROM CS_TABLE
```

```
WHERE PREV = "Hotel_California" AND
```

```
TYPE = "link"
```

```
ORDER BY PREV DESC;
```



### 3. What series of Wikipedia articles, starting with Hotel California, keeps the largest fraction of its readers clicking on internal links?

#### HIVE QUERIES USED TO ANSWER THE QUESTION (cont)


Now that we have the data that contains only internal links, we must aggregate the num of internal link clicks with similar curr page\_titles and insert new data into a new table

```
3) CREATE TABLE Q3_RESULTS (PREV STRING, CURR STRING, NUM_OF_CLICK_THROUGHS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

```
4) INSERT INTO Q3_RESULTS (SELECT PREV, CURR, SUM (NUM_OF_CLICK_THROUGHS) AS NUM_OF_INTERNAL  
FROM INTERNAL_DATA2 GROUP BY CURR);
```

Display the results of the table and order them by the num of internal clicks

```
5) SELECT * FROM Q3_RESULTS ORDER BY NUM_OF_INTERNAL LIMIT 10;
```




4. Find an example of an English Wikipedia article that is relatively more popular in the UK. Find the same for the US and Australia.

Assumptions

- 1) According to the latest study of global internet traffic by Cisco, the internet is busiest between 9pm and 11pm around the world.
- 2) Wiki page views traffic time stamps are based on UTC (Coordinated Universal Time)
- 3) So if UTC is 12:00:00am, time in US eastern is 07:00:00pm, time in US pacific 04:00:00pm, time in US central is 06:00:00pm, time in UK (UTC) is 12:03:00am, and time in Australia (GMT) is 11:00:00am.

US eastern 9PM-11PM	UTC 2AM-4AM
US central 9PM-11PM	UTC 3AM - 5AM
US pacific 9PM-11PM	UTC 5AM - 7AM
UK 9PM-11PM	UTC 9PM - 11PM
Australia 9PM-11PM	UTC 8AM - 10AM



#### 4. Find an example of an English Wikipedia article that is relatively more popular in the UK. Find the same for the US and Australia.


1) We can look at the time posted in <https://dumps.wikimedia.org/other/pageviews/2020/2020-11/>

2) Assuming that Wiki is running on UTC (Coordinated Universal Time), we can take a .gz from (for example: 06-Nov-2020 01:43) and assume that the English articles contained inside that .gz are relatively more popular in the US.

If from (05-Nov-2020 20:49), the English articles contained in that .gz are more popular in the UK

If from (05-Nov-2020 08:49), the English articles contained in that .gz are more popular in Australia

<a href="#">pageviews-20201105-050000.gz</a>	05-Nov-2020 05:45	37413060
<a href="#">pageviews-20201105-060000.gz</a>	05-Nov-2020 06:43	39301290
<a href="#">pageviews-20201105-070000.gz</a>	05-Nov-2020 07:50	43496992
<a href="#">pageviews-20201105-080000.gz</a>	05-Nov-2020 08:49	49093413
<a href="#">pageviews-20201105-090000.gz</a>	05-Nov-2020 09:49	52479979
<a href="#">pageviews-20201105-100000.gz</a>	05-Nov-2020 10:48	54514071
<a href="#">pageviews-20201105-110000.gz</a>	05-Nov-2020 11:50	56142654
<a href="#">pageviews-20201105-120000.gz</a>	05-Nov-2020 12:52	57467962
<a href="#">pageviews-20201105-130000.gz</a>	05-Nov-2020 13:52	59480695
<a href="#">pageviews-20201105-140000.gz</a>	05-Nov-2020 14:54	61862080
<a href="#">pageviews-20201105-150000.gz</a>	05-Nov-2020 15:50	62854280
<a href="#">pageviews-20201105-160000.gz</a>	05-Nov-2020 16:50	62412700
<a href="#">pageviews-20201105-170000.gz</a>	05-Nov-2020 18:01	61859517
<a href="#">pageviews-20201105-180000.gz</a>	05-Nov-2020 18:51	60550577
<a href="#">pageviews-20201105-190000.gz</a>	05-Nov-2020 19:51	59618348
<a href="#">pageviews-20201105-200000.gz</a>	05-Nov-2020 20:49	58980558
<a href="#">pageviews-20201105-210000.gz</a>	05-Nov-2020 21:49	58279356
<a href="#">pageviews-20201105-220000.gz</a>	05-Nov-2020 22:47	55404039
<a href="#">pageviews-20201105-230000.gz</a>	05-Nov-2020 23:49	50827772
<a href="#">pageviews-20201106-000000.gz</a>	06-Nov-2020 00:50	44968798
<a href="#">pageviews-20201106-010000.gz</a>	06-Nov-2020 01:43	40880515
<a href="#">pageviews-20201106-020000.gz</a>	06-Nov-2020 02:41	39041593
<a href="#">pageviews-20201106-030000.gz</a>	06-Nov-2020 03:49	38626148
<a href="#">pageviews-20201106-040000.gz</a>	06-Nov-2020 04:43	38807947



5. Analyze how many users will see the average vandalized Wikipedia page before the offending edit is reversed.

### SOURCES USED

We will be using September's data


pageviews-20200901-000000.gz - pageviews-20200930-230000.gz

2020-09.enwiki.2020-09.tsv.bz2

### HIVE QUERIES USED TO ANSWER THE QUESTION

Create a table that can hold all of the data in "2020-09.enwiki.2020-09.tsv.bz2."

```
1) CREATE TABLE DUMP_DATA (  
  
    wiki_db string,  
  
    event_entity string,  
  
    event_type string,  
  
    ...  
  
    revision_tags array<string>  
  
) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```



## 5. Analyze how many users will see the average vandalized Wikipedia page before the offending edit is reversed.

### HIVE QUERIES USED TO ANSWER THE QUESTION (cont)


Create tables to insert data and group by

```
2) CREATE TABLE FILTERED_DATA (PAGE_TITLE STRING, revision_seconds_to_identity_revert BIGINT) ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '\t';
```

```
.3) INSERT INTO FILTERED_DATA (SELECT PAGE_TITLE, revision_seconds_to_identity_revert FROM DUMP_DATA WHERE  
EVENT_ENTITY='revision' AND revision_seconds_to_identity_revert > 0.0);
```

```
4) CREATE TABLE UPDATED_FILTERED_DATA (PAGE_TITLE STRING, revision_seconds_to_identity_revert BIGINT) ROW FORMAT  
DELIMITED FIELDS TERMINATED BY '\t';
```

```
5) INSERT INTO UPDATED_FILTERED_DATA (SELECT PAGE_TITLE, SUM (revision_seconds_to_identity_revert) FROM FILTERED_DATA  
GROUP BY PAGE_TITLE);
```



## 5. Analyze how many users will see the average vandalized Wikipedia page before the offending edit is reversed.

### HIVE QUERIES USED TO ANSWER THE QUESTION (cont)

Reuse the `UPDATED_EN_PAGEVIEWS` table from earlier that contains the pageview data from Sept with domain\_code "en" and "en.m" and inner join the table with `UPDATED_FILTERED_DATA` (return the rows that have matching page\_title and move the results to `Q5_RESULTS`)

```
6) CREATE TABLE Q5_RESULTS AS
```

```
SELECT UPDATED_EN_PAGEVIEWS.PAGE_TITLE, UPDATED_EN_PAGEVIEWS.NUM_OF_PAGE_VIEWS,  
UPDATED_FILTERED_DATA.revision_seconds_to_identity_revert
```

```
FROM UPDATED_EN_PAGEVIEWS
```

```
INNER JOIN UPDATED_FILTERED_DATA ON UPDATED_EN_PAGEVIEWS.PAGE_TITLE = UPDATED_FILTERED_DATA.PAGE_TITLE;
```

Now we want to make our computation:  $(\text{SUM}(\text{NUM\_OF\_PAGE\_VIEWS}) / \text{COUNT}(\text{PAGE\_TITLE}))$  FROM `UPDATED_FILTERED_DATA` to find the avg number of users that saw the vandalized Wiki page before the edit is reversed.

```
7) SELECT SUM(NUM_OF_PAGE_VIEWS) / COUNT(PAGE_TITLE) FROM Q5_RESULTS;
```

## 6. What is the total number of English articles visited on 10/20/2020 with titles that start with the letter 'a'?

### SOURCES USED

pageviews-20201020-000000.gz - pageviews-20201020-230000.gz

### HIVE QUERIES USED TO ANSWER THE QUESTION

Similar procedure as question 1. First we create a table that takes in pageview data from Oct containing only domain codes 'en' and 'en.m'.

1) USE PAGEVIEW\_DB;

2) CREATE TABLE OCT\_PAGEVIEWS (DOMAIN\_CODE STRING, PAGE\_TITLE STRING, NUM\_OF\_REQUESTS INT, RESPONSE\_BYTES INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ' ';


3) CREATE TABLE EN\_PAGEVIEWS (PAGE\_TITLE STRING, NUM\_OF\_REQUESTS INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ' ';

4) LOAD DATA LOCAL INPATH '/home/aliang30/pageviews-20201020/\*' INTO TABLE OCT\_PAGEVIEWS;

5) INSERT INTO EN\_PAGEVIEWS (SELECT PAGE\_TITLE, NUM\_OF\_REQUESTS FROM OCT\_PAGEVIEWS WHERE DOMAIN\_CODE='en');

6) INSERT INTO EN\_PAGEVIEWS (SELECT PAGE\_TITLE, NUM\_OF\_REQUESTS FROM OCT\_PAGEVIEWS WHERE DOMAIN\_CODE='en.m');





6. What is the total number of English articles visited on 10/20/2020 with titles that start with the letter 'a'?

HIVE QUERIES USED TO ANSWER THE QUESTION (cont)

Now that we have our table with only English wiki articles, we can find the total number of articles that start with the letter 'a'

```
7) SELECT COUNT(*) AS total_a  
FROM EN_PAGEVIEWS  
WHERE PAGE_TITLE LIKE 'A%';
```

```
OK  
+-----+  
| total_a |  
+-----+  
| 3347674 |  
+-----+  
1 row selected (33.562 seconds)  
0: jdbc:hive2://> |
```



## Github link

[https://github.com/AlanLiang1/Project\\_1.git](https://github.com/AlanLiang1/Project_1.git)